

Sparse Node Perturbation for FPGA-Oriented Online Learning

Progress Report — does sparse (and *adaptive* sparse) really work?

Melike Kara

June 2026

The question we’re actually answering. *Can adaptive sparse node perturbation reach the same convergence and accuracy as full node perturbation, while cutting the training energy and compute cost in an FPGA-based online-learning system?*

The short answer. Sparse perturbation works: perturbing a small fraction of neurons keeps almost all of the accuracy at a fraction of the cost. But the *adaptive* part — trying to be clever about *which* neurons to perturb — does **not** beat simply spreading the perturbation around (random or round-robin). The reason is concrete: the learning signal is spread almost evenly across neurons, so there are no “important” neurons to pick. The practical takeaway is good news for hardware: the cost savings come from cheap, sort-free sparse coverage, not from expensive ranking logic. On top of that, replacing Gaussian noise with single-bit (Rademacher) noise costs nothing in accuracy.

All numbers below are read from **results/**; perturbation experiments use 3 seeds. Treat differences smaller than the standard deviation as noise.

1 How it works, and the two experiment families

(D)ANP in one paragraph. Add a little random noise to the neuron activations, see how much the loss changes, and nudge the weights in the direction that helped:

$$\Delta W_l = N \delta L \frac{\delta a_l}{\|\delta a\|^2} x_{l-1}^\top,$$

where δL is the loss change between the clean and noisy pass and δa_l is how much the activities moved. No backward pass — just two forward passes. DANP adds a decorrelation step.

Why $\sigma = 0.01$? The noise scale isn’t a free tuning knob — the math pins it to a band. Too small and δL (a difference of two losses) sinks below floating-point rounding, so you measure numerical junk instead of a real loss change (the reference paper sees this below $\sigma^2 \approx 10^{-7}$). Too large and the first-order approximation $\delta L \approx \nabla L \cdot v$ breaks — the perturbation leaves the linear region and the curvature term distorts the estimate. That leaves a usable band of roughly $\sigma^2 \in [10^{-6}, 10^{-3}]$. We picked $\sigma = 0.01$ ($\sigma^2 = 10^{-4}$), the middle of that band: big enough to measure cleanly, small enough to stay linear. The reference uses 10^{-6} ; ours is two orders larger but still in the safe band, and — importantly — it is held *fixed* across every run, so it is a constant, not a variable in any comparison.

Two families of experiments. We have results from two setups; keep them separate when reading the tables:

- **DANP, wide net** [1024, 1024, 1024] (3082 neurons): the main test of sparse-vs-full and of *adaptive* selection (top-k policies), plus the gradient-distribution analysis and the width sweep. MNIST and CIFAR-10.
- **ANP, small net** [64, 64] (138 neurons): the budget sweep down to extreme sparsity, Gaussian-vs-Rademacher noise, and the antithetic test. MNIST and Fashion-MNIST.

Policies. *Random* = a fresh random set of neurons each step. *Scheduled* = round-robin, every neuron visited once per H/k steps (just a counter in hardware). *Adaptive top-k* = pick the k “best” neurons by a score: activity-difference, activity-loss contribution, or gradient-alignment. Optimizer Adam, cross-entropy loss.

2 Does sparse match full? (yes)

On the wide DANP network, dropping to 25% active neurons barely moves accuracy — on MNIST it stays within about 0.8% of full, and on CIFAR-10 it’s essentially unchanged (Tables 1, 2). That’s a $4\times$ **cut in perturbation/update cost** for a sub-1% accuracy hit.

Table 1: DANP on MNIST, [1024³], 3 seeds. “Cost” = active/total neurons. **Note the epoch mismatch:** random/scheduled 25% ran 20 epochs, but the adaptive top-k runs ran only 5 — so the fair adaptive-vs-simple comparison is at 50%/5-epochs (all ≈ 0.94), not 25%.

Policy	Fraction	Cost	Final acc	Epochs
Full DANP	—	1.000	0.9642	20
Random	25%	0.250	0.9558	20
Scheduled	25%	0.250	0.9566	20
Random	50%	0.500	0.9407	5
ActivityDiff top-k	50%	0.500	0.9377	5
ActivityLoss top-k	50%	0.500	0.9398	5
ActivityDiff top-k	25%	0.250	0.9313	5
ActivityLoss top-k	25%	0.250	0.9314	5

Table 2: DANP on CIFAR-10, [1024³], 5 epochs, 3 seeds. All policies ran the same 5 epochs, so this is the clean comparison. An MLP tops out around 0.41 on CIFAR, so these are all near the ceiling for this model — and they’re all the same.

Policy	Fraction	Cost	Final acc
Full DANP	—	1.000	0.4094
Random	50%	0.500	0.4080
Random	25%	0.250	0.4070
Scheduled	50%	0.500	0.4122
Scheduled	25%	0.250	0.4102
ActivityLoss top-k	50%	0.500	0.4109
ActivityLoss top-k	25%	0.250	0.4055
GradientAligned top-k	50%	0.500	0.4117
GradientAligned top-k	25%	0.250	0.4076

(Side note: this DANP CIFAR run *converges* at ~ 0.41 . An earlier *plain-ANP* CIFAR pilot on a different, smaller net was unstable and collapsed — so decorrelation is what makes CIFAR behave. See §9.)

3 Does “adaptive” (smart selection) help? (no)

This is the heart of the research question. The adaptive top-k policies pick neurons by a cleverness score instead of just covering everyone. The result: at matched settings they do **not** beat random or scheduled. On CIFAR (all 5 epochs, fair comparison) every policy lands at ≈ 0.41 — adaptive, random, scheduled, all the same. On MNIST at matched epochs (50%, 5 epochs) it’s the same story: random 0.9407 vs. activity-diff 0.9377 vs. activity-loss 0.9398 — a wash. The low MNIST 25% top-k numbers (0.931) look worse only because those runs got 5 epochs instead of 20; don’t read them as a real gap.

So the expensive part — ranking/sorting neurons every step — buys nothing.

4 Why adaptive doesn’t help: the signal is diffuse

We measured how concentrated the per-neuron learning signal $g_i \propto \Delta a_i \Delta L / \|\Delta \mathbf{a}\|^2$ is (Table 3). Three numbers say the same thing: the coefficient of variation stays around 0.25–0.28, the Gini coefficient around 0.13–0.16, and the top 10% of neurons carry only ~ 13 –15% of the total signal energy — barely above the 10% you’d get if every neuron were identical.

In plain terms: **no small group of neurons is special**. The useful signal is smeared almost evenly across all of them. If there’s nothing special to find, then “smartly selecting the important neurons” has nothing to exploit — which is exactly why the top-k policies tie with random. The thing that *does* matter is making sure every neuron gets its turn over time, i.e. **coverage**, and a round-robin counter already gives that for free.

Table 3: How concentrated the per-neuron signal is, full DANP [1024³], 5 epochs, mean over 3 seeds. Low CV / low Gini / low top-10% share = diffuse, i.e. no neurons stand out. Stays diffuse throughout training on both datasets.

Epoch	MNIST			CIFAR-10		
	CV	Top10%	Gini	CV	Top10%	Gini
1	0.238	0.135	0.134	0.228	0.132	0.129
2	0.253	0.138	0.142	0.234	0.133	0.132
3	0.263	0.141	0.148	0.238	0.134	0.134
4	0.272	0.143	0.152	0.241	0.134	0.136
5	0.281	0.144	0.156	0.245	0.135	0.138

5 Does it hold across network sizes? (yes)

Repeating scheduled-25% across three widths, the accuracy drop stays small everywhere (Table 4). Note the wider net even improves in absolute accuracy — DANP scales up nicely. (Plain ANP, by contrast, gets *worse* with width because the wide layer hogs the global $\|\delta \mathbf{a}\|^2$ normaliser; that’s one reason DANP is the right base for the big nets.)

Table 4: Width sweep, MNIST, scheduled 25%, 20 epochs, 3 seeds.

Architecture	Cost	Full acc	Sparse acc	Drop
[64, 64, 64]	0.248	0.9391	0.9342	−0.0049
[256, 256, 256]	0.249	0.9530	0.9480	−0.0050
[1024, 1024, 1024]	0.250	0.9642	0.9566	−0.0076

6 Random vs. scheduled, and cheap noise

Random vs. scheduled. On the small net they’re basically tied across fractions (Table 6) — so *how many* neurons matters, *which ones* mostly doesn’t. Scheduled still wins on hardware (a counter, no random bits to build the mask). The one place the policy clearly mattered was the wide $H=1024$ MNIST run: random lost 5.9% vs. full, scheduled only 1.5% — so coverage control starts to matter when the per-step signal weakens with width.

Rademacher = Gaussian. Swapping multi-bit Gaussian noise for $\pm\sigma$ single-bit (Rademacher) noise changes accuracy by less than the seed-to-seed spread (Table 5). Same learning, single bit flip instead of a Gaussian sampler — a clean hardware win at zero accuracy cost.

Table 5: MNIST, ANP [64, 64], scheduled: Gaussian vs. Rademacher (best acc $\pm$ std, $n = 3$). Every difference is within the std.

α	Gaussian	Rademacher	diff
0.0625	0.9257 ± 0.0026	0.9258 ± 0.0023	+0.0001
0.125	0.9266 ± 0.0028	0.9288 ± 0.0005	+0.0022
0.25	0.9262 ± 0.0008	0.9277 ± 0.0028	+0.0015
0.5	0.9299 ± 0.0019	0.9280 ± 0.0013	−0.0019

7 How sparse can we go before it breaks?

MNIST is too easy to show a limit (flat curve, Fig. 1b). Fashion-MNIST shows a real one (Table 6, Fig. 1a): accuracy climbs, then flattens with a clear knee around 12–25%. At 25% we’re at 0.8248, indistinguishable from dense (0.8191); at just 12% we’re within 0.6% of the half-network run — so $\sim 88\%$ of the perturbation budget is genuinely wasted.

It only breaks at the very bottom. At $\alpha = 0.016$ (≈ 3 neurons) one seed trains fine then collapses late (Fig. 2). Not a bug: with 3 active neurons, each is revisited only every ~ 46 steps, so coverage gaps let the estimate drift. That’s the **coverage floor** — the same lesson as §4, now from the other direction.

Table 6: Fashion-MNIST, ANP [64, 64], Rademacher (best acc $\pm$ std, $n = 3$). Dense baseline 0.8191 ± 0.0036 .

α	Cost	Scheduled	Random
0.016	0.022	0.7863 ± 0.0075	0.7903 ± 0.0045
0.031	0.036	0.8043 ± 0.0017	0.8102 ± 0.0030
0.063	0.065	0.8114 ± 0.0052	0.8140 ± 0.0049
0.125	0.123	0.8147 ± 0.0079	0.8137 ± 0.0040
0.25	0.246	0.8248 ± 0.0035	0.8161 ± 0.0014
0.5	0.500	0.8206 ± 0.0044	0.8228 ± 0.0033

8 One trick that didn’t pay off: antithetic sampling

Idea. Use the same noise as $+v$ and $-v$ and average the two estimates; in theory this cancels the curvature error in δL and lowers variance for one extra pass. **Result.** It barely moved (Table 7): antithetic $\approx$ plain single-pass, and the fair 2-pass control (“resample”, two *independent* samples) beat it at every fraction. **Why.** Our noise is small, so the curvature error antithetic removes is already negligible — it cancels something that barely matters. The real noise at extreme sparsity is first-order sampling variance, which only independent averaging touches. And since both 2-pass




Figure 1: Best test accuracy vs. how many neurons we perturb (± 1 std, $n = 3$), ANP [64, 64]. (a) Fashion-MNIST: a real curve with a knee around 12–25%. (b) MNIST: flat — too easy to stress the method.




Figure 2: Below the coverage floor: with ≈ 3 active neurons, seed 43 trains fine then collapses at the end. Each neuron is touched only every ~ 46 steps.

modes *double* the perturbation cost — the opposite of our goal — it’s not worth it. We keep single-pass sparse Rademacher.

Table 7: Fashion-MNIST, scheduled Rademacher, extreme regime (best acc $\pm$ std, $n = 3$). Resample beats antithetic everywhere; both cost 2 passes.

α	single (1 pass)	antithetic (2)	resample (2)
0.016	0.7935 ± 0.0038	0.7945 ± 0.0070	0.7987 ± 0.0048
0.031	0.8068 ± 0.0030	0.8046 ± 0.0031	0.8132 ± 0.0032
0.063	0.8076 ± 0.0021	0.8124 ± 0.0037	0.8165 ± 0.0025

9 CIFAR-10: DANP converges, plain ANP didn’t

The DANP CIFAR runs above are stable at ~ 0.41 (the ceiling for an MLP). An earlier *plain-ANP* pilot on a smaller $[256, 256]$ net was unstable — loss blew up and runs fell to chance — so we can’t draw conclusions from that pilot. The contrast is itself informative: decorrelation is what keeps CIFAR stable. Fixing the plain-ANP pilot (input standardisation, an ϵ in the $\|\delta a\|^2$ denominator, update clipping) is a next step, not a result.

10 Bottom line — answering the question

- **Sparse matches full.** 25% active neurons keeps accuracy within $<1\%$ on MNIST and essentially unchanged on CIFAR, at a $4\times$ cost cut.
- **Adaptive selection does *not* help.** Top-k by any score ties with random/scheduled, because the signal is diffuse (no special neurons). So the answer to “does *adaptive* sparse perturbation help?” is: not on these tasks — and that’s good for FPGA, since it means no ranking/sorting hardware.
- **Coverage is the real lever.** Spreading perturbation over all neurons over time is what matters; a round-robin counter delivers it for free, and the method breaks only below a small coverage floor.
- **Cheaper still:** Rademacher (single-bit) noise matches Gaussian, and the 2-pass variance tricks (antithetic/resample) aren’t worth their doubled cost.

11 What’s next

1. **Stabilise plain-ANP CIFAR** (standardise inputs, ϵ in the denominator, clipping).
2. **FPGA cost model:** area/energy of a single-bit Rademacher source vs. a Gaussian sampler, and α mapped to a power budget.